

# PLUG AND PLAY AUTOENCODERS FOR CONDITIONAL TEXT GENERATION

FLORIAN MAI †, ♠ NIKOLAOS PAPPAS ♣

IVAN MONTERO ♣ NOAH A. SMITH ♣, ◇ JAMES HENDERSON †

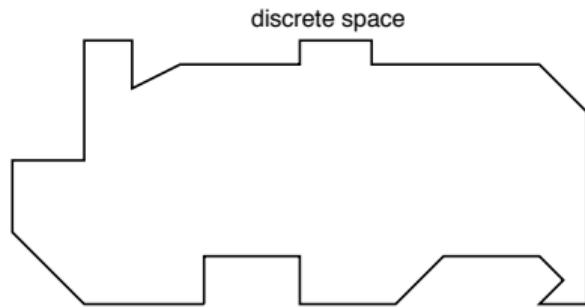
† IDIAP RESEARCH INSTITUTE, ♠ EPFL, SWITZERLAND

♣ UNIVERSITY OF WASHINGTON, SEATTLE, USA

◇ ALLEN INSTITUTE FOR ARTIFICIAL INTELLIGENCE, SEATTLE, USA

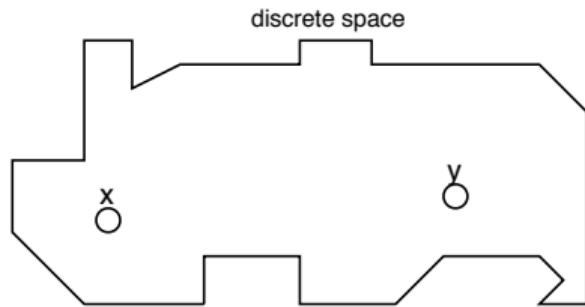


# THE PROBLEM WITH CONDITIONAL TEXT GENERATION



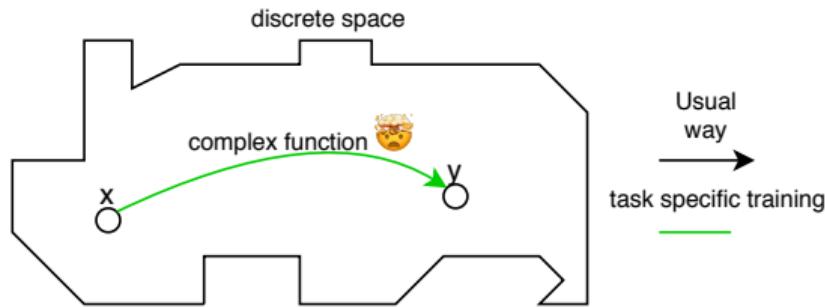
- text lives in a messy, discrete space

# THE PROBLEM WITH CONDITIONAL TEXT GENERATION



- text lives in a messy, discrete space
- conditional text generation requires mapping from discrete input to discrete output

# THE PROBLEM WITH CONDITIONAL TEXT GENERATION

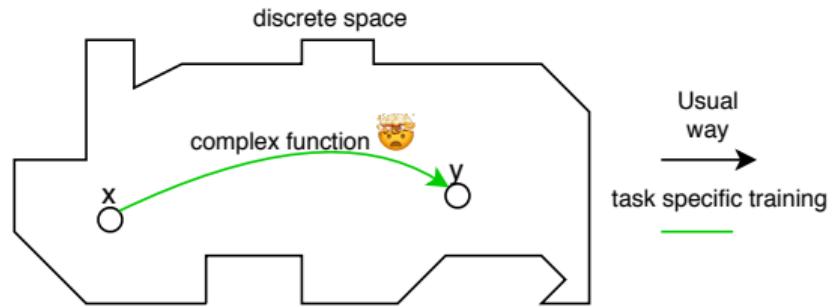


- text lives in a messy, discrete space
- conditional text generation requires mapping from discrete input to discrete output

## Usual way:

- learning a complex, task-specific function, which is difficult to train in discrete space

# THE PROBLEM WITH CONDITIONAL TEXT GENERATION



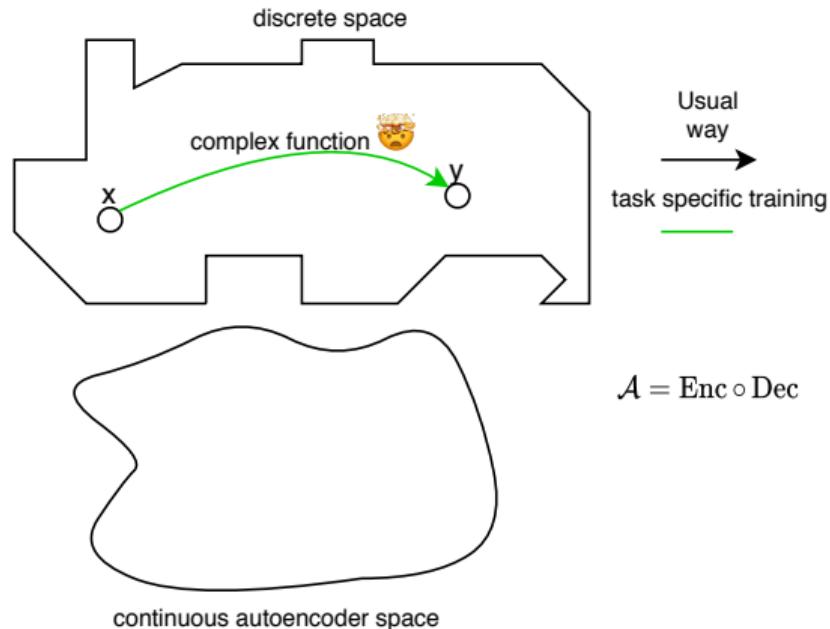
- text lives in a messy, discrete space
- conditional text generation requires mapping from discrete input to discrete output

## Usual way:

- learning a complex, task-specific function, which is difficult to train in discrete space

## Our way:

# THE PROBLEM WITH CONDITIONAL TEXT GENERATION



- text lives in a messy, discrete space
- conditional text generation requires mapping from discrete input to discrete output

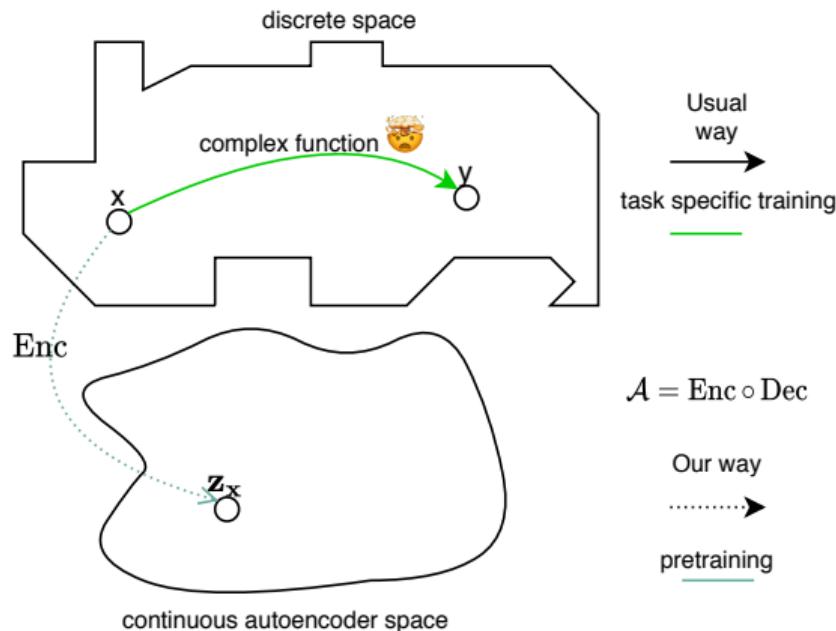
## Usual way:

- learning a complex, task-specific function, which is difficult to train in discrete space

## Our way:

- obtain a continuous space by training an autoencoder

# THE PROBLEM WITH CONDITIONAL TEXT GENERATION



- text lives in a messy, discrete space
- conditional text generation requires mapping from discrete input to discrete output

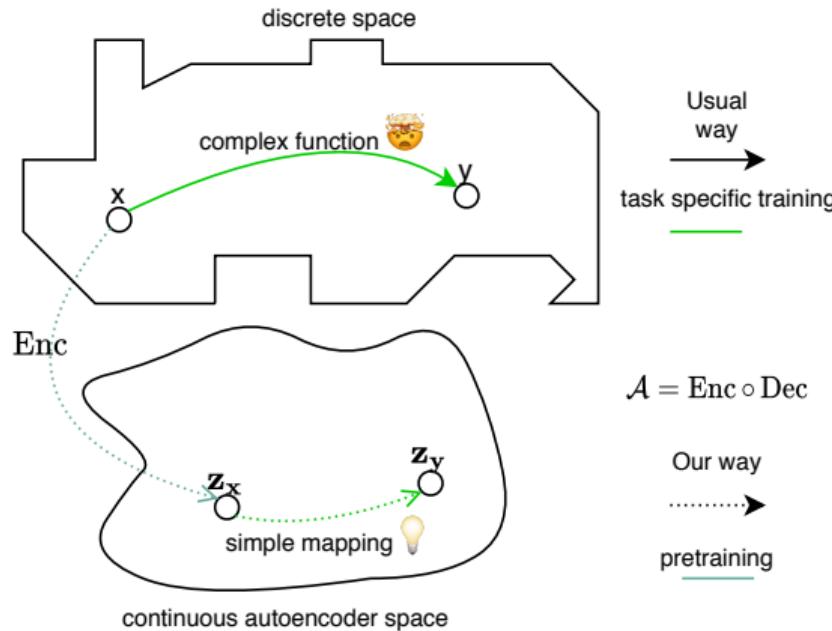
## Usual way:

- learning a complex, task-specific function, which is difficult to train in discrete space

## Our way:

- obtain a continuous space by training an autoencoder

# THE PROBLEM WITH CONDITIONAL TEXT GENERATION



- text lives in a messy, discrete space
- conditional text generation requires mapping from discrete input to discrete output

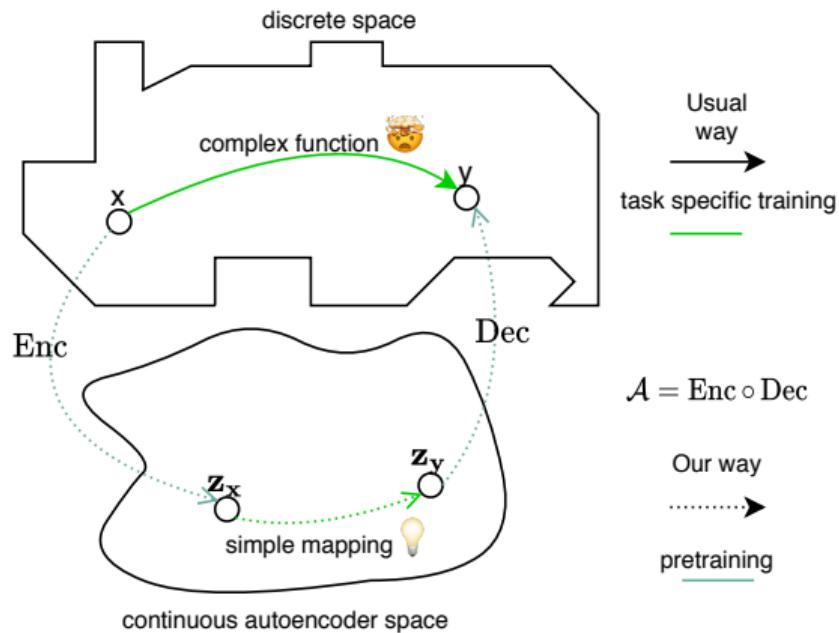
## Usual way:

- learning a complex, task-specific function, which is difficult to train in discrete space

## Our way:

- obtain a continuous space by training an autoencoder
- reduce task-specific learning to the continuous space

# THE PROBLEM WITH CONDITIONAL TEXT GENERATION



- text lives in a messy, discrete space
- conditional text generation requires mapping from discrete input to discrete output

## Usual way:

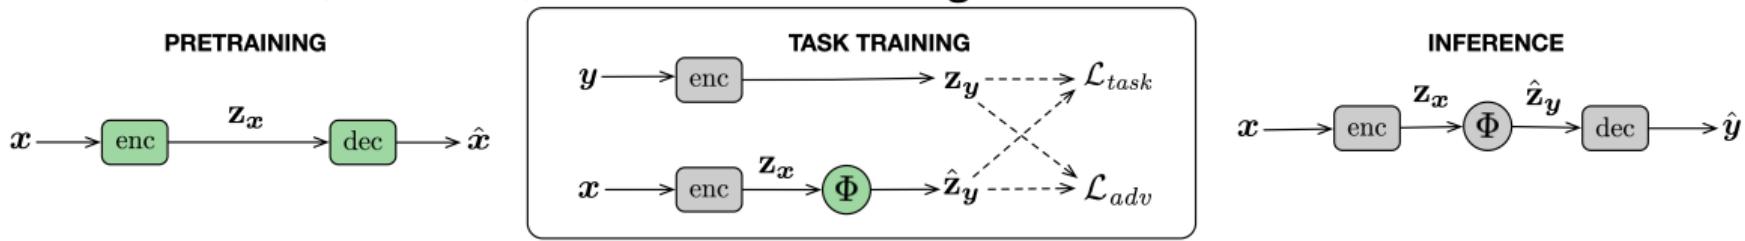
- learning a complex, task-specific function, which is difficult to train in discrete space

## Our way:

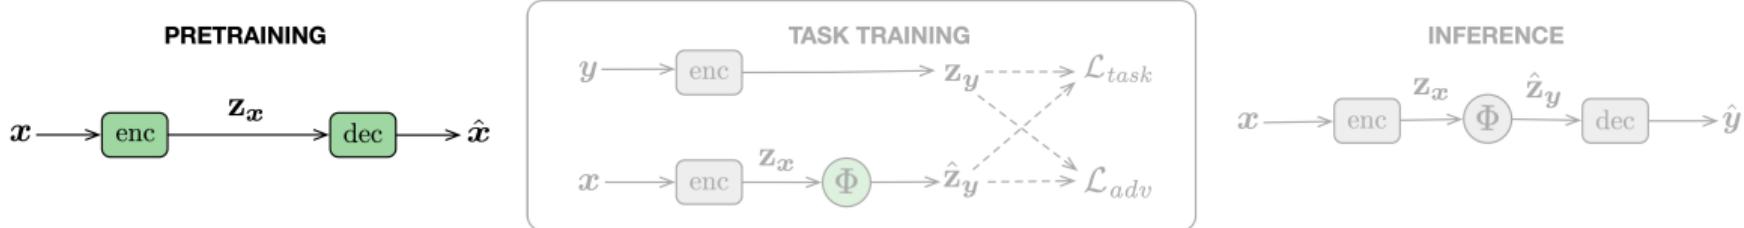
- obtain a continuous space by training an autoencoder
- reduce task-specific learning to the continuous space

# FRAMEWORK OVERVIEW

Our framework (*Emb2Emb*) consists of three stages:



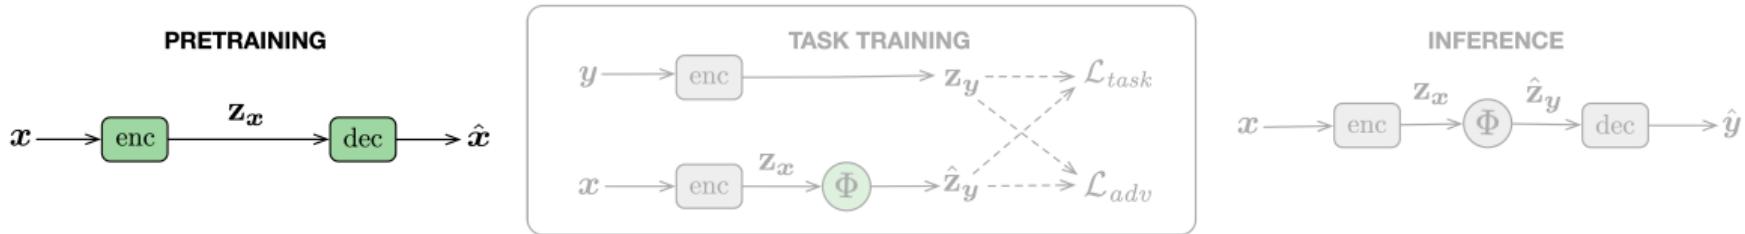
# FRAMEWORK OVERVIEW



## Pretraining:

- Train a model of the form  $\mathcal{A}(x) = \text{Dec}(\text{Enc}(x))$  on corpus of sentences
- Assume a fixed-size continuous embedding  $\mathbf{z}_x := \text{Enc}(x) \in \mathbb{R}^d$
- Enc and Dec can be any function trained with any objective so long as  $\mathcal{A}(x) \approx x$
- training corpus can be any unlabeled corpus  $\Rightarrow$  large-scale pretraining?

# FRAMEWORK OVERVIEW



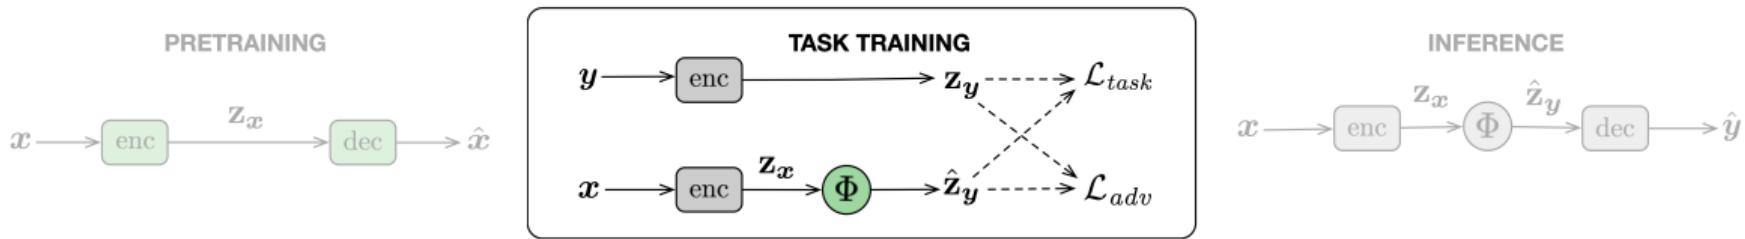
## Pretraining:

- Train a model of the form  $\mathcal{A}(x) = \text{Dec}(\text{Enc}(x))$  on corpus of sentences
- Assume a fixed-size continuous embedding  $\mathbf{z}_x := \text{Enc}(x) \in \mathbb{R}^d$
- Enc and Dec can be any function trained with any objective so long as  $\mathcal{A}(x) \approx x$
- training corpus can be any unlabeled corpus  $\Rightarrow$  large-scale pretraining?

## Plug and Play

Our framework is *plug and play* because any autoencoder can be used with it.

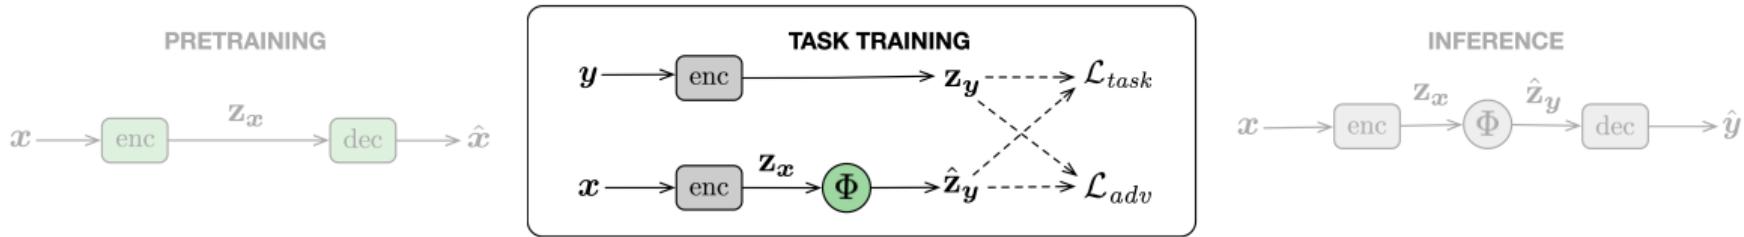
# FRAMEWORK OVERVIEW



## Task Training:

- Supervised case:  $\mathcal{L}_{task}(\hat{\mathbf{z}}_y, \mathbf{z}_y) = d(\hat{\mathbf{z}}_y, \mathbf{z}_y)$  where  $d$  is a distance function (cosine distance loss in our experiments).

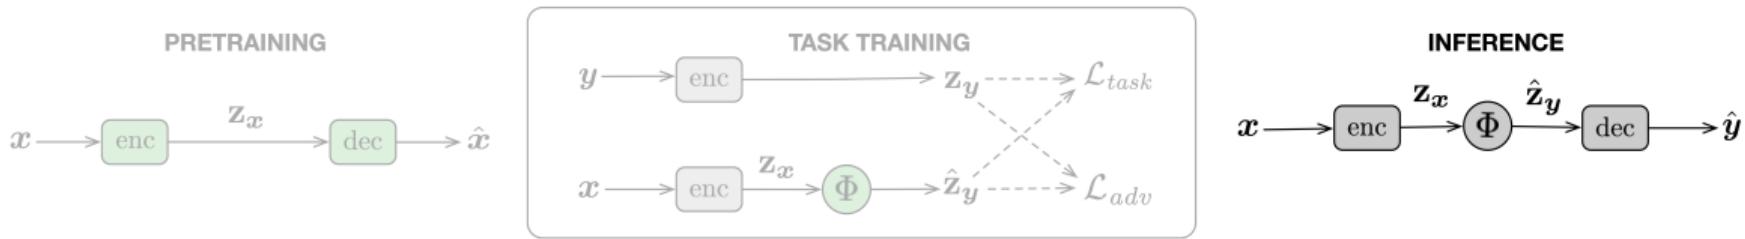
# FRAMEWORK OVERVIEW



## Task Training:

- Supervised case:  $\mathcal{L}_{task}(\hat{\mathbf{z}}_y, \mathbf{z}_y) = d(\hat{\mathbf{z}}_y, \mathbf{z}_y)$  where  $d$  is a distance function (cosine distance loss in our experiments).
- Training objective:  $\mathcal{L} = \mathcal{L}_{task} + \lambda_{adv} \cdot \boxed{\mathcal{L}_{adv}}$  ?

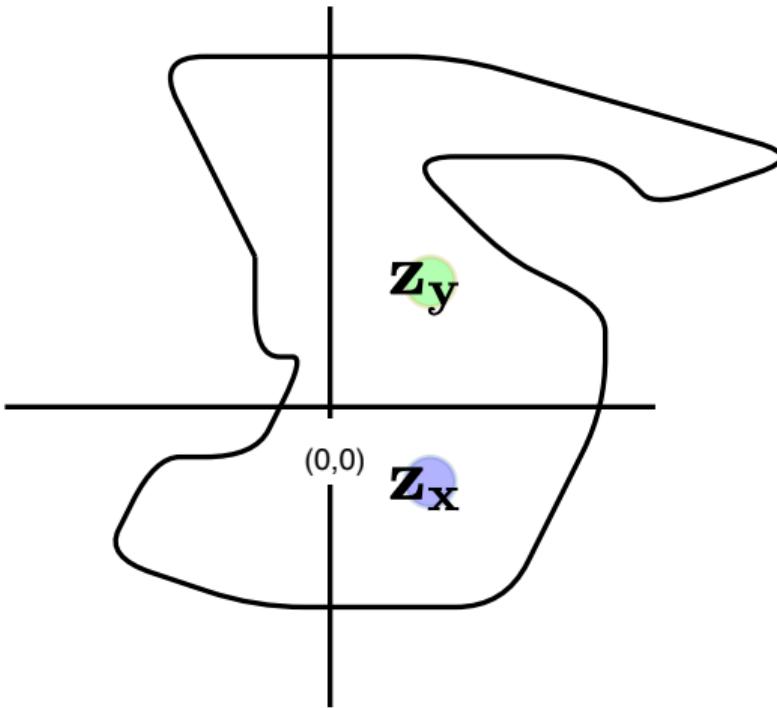
# FRAMEWORK OVERVIEW



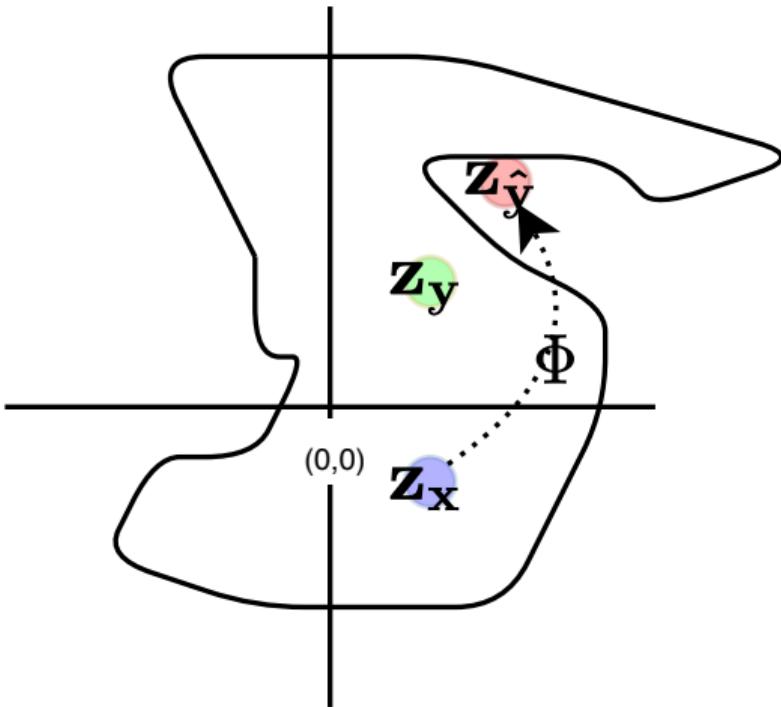
## Inference:

- compose inference model as  $\text{Enc} \circ \Phi \circ \text{Dec}$
- but: Dec not involved in training. Can it handle outputs of  $\Phi$ ?
- $\Rightarrow$  yes, if using  $\mathcal{L}_{adv}$ .

## WHAT CAN HAPPEN WHEN LEARNING IN THE EMBEDDING SPACE?

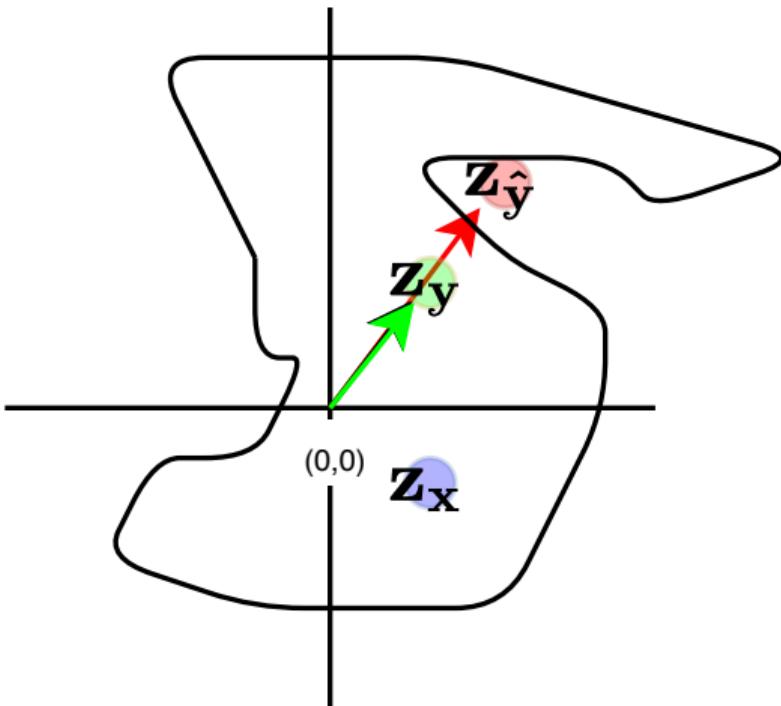


## WHAT CAN HAPPEN WHEN LEARNING IN THE EMBEDDING SPACE?



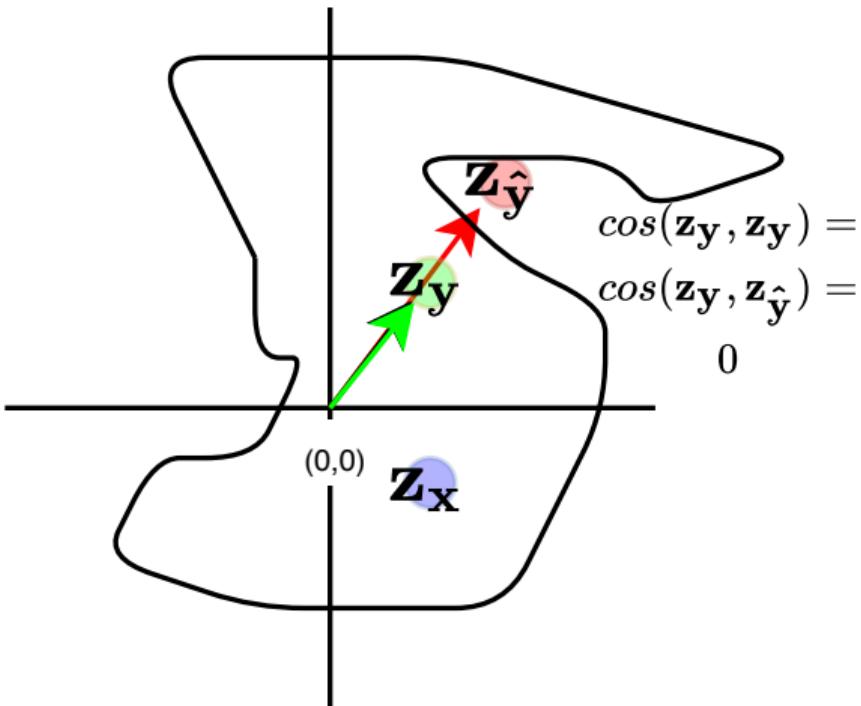
- A prediction may end up **off** the manifold, and by definition, the decoder cannot handle off-manifold data well, but ...

## WHAT CAN HAPPEN WHEN LEARNING IN THE EMBEDDING SPACE?



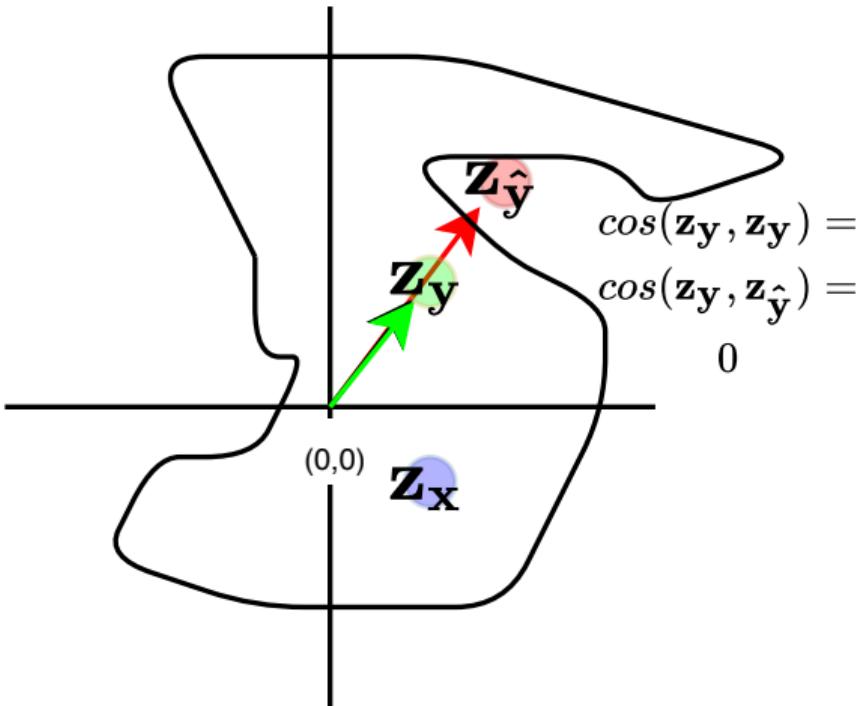
- A prediction may end up **off** the manifold, and by definition, the decoder cannot handle off-manifold data well, but ...
- ... but the predicted embedding may still have the same angle as the true output embedding...

## WHAT CAN HAPPEN WHEN LEARNING IN THE EMBEDDING SPACE?



- A prediction may end up **off** the manifold, and by definition, the decoder cannot handle off-manifold data well, but ...
- ... but the predicted embedding may still have the same angle as the true output embedding...
- resulting in zero cosine distance loss despite being off the manifold.

## WHAT CAN HAPPEN WHEN LEARNING IN THE EMBEDDING SPACE?



- A prediction may end up **off** the manifold, and by definition, the decoder cannot handle off-manifold data well, but ...
- ... but the predicted embedding may still have the same angle as the true output embedding...
- resulting in zero cosine distance loss despite being off the manifold.
- Similar problems arise for L<sub>2</sub> distance - how do we keep the embeddings **on** the manifold?

## ADVERSARIAL LOSS TERM

- train a discriminator  $\text{disc}$  to distinguish between embeddings produced by the encoder and embeddings resulting from the mapping:

$$\max_{\text{disc}} \sum_{i=1}^N \log(\text{disc}(\mathbf{z}_{\tilde{\mathbf{y}}_i})) + \log(1 - \text{disc}(\Phi(\mathbf{z}_{\mathbf{x}_i}))$$

- using the adversarial learning framework, mapping acts as the adversary and tries to fool the discriminator:

$$\mathcal{L}_{adv}(\Phi(\mathbf{z}_{\mathbf{x}_i}); \theta) = -\log(\text{disc}(\Phi(\mathbf{z}_{\mathbf{x}_i}); \theta))$$

- at convergence, the mapping should only produce embeddings that are on the manifold

## SUPERVISED STYLE TRANSFER EXPERIMENTS

- *WikiLarge* dataset: transform “normal” English to “simple” English
- parallel sentences (input and output) are available

| Model                             | BLEU (relative imp.) | SARI (relative imp.) |
|-----------------------------------|----------------------|----------------------|
| Emb2Emb (no $\mathcal{L}_{adv}$ ) | 15.7 (-)             | 21.1 (-)             |
| Emb2Emb                           | <b>34.7</b> (+121%)  | <b>25.4</b> (+20.4%) |

The adversarial loss term  $\mathcal{L}_{adv}$  is crucial for embedding-to-embedding training!

## SUPERVISED STYLE TRANSFER EXPERIMENTS

- we conducted controlled experiments of models **with a fixed-size bottleneck**
- best Seq2Seq model: best performing variant among fixed-size bottleneck models that are trained end-to-end via token-level cross-entropy loss (like Seq2Seq)

| Model              | BLEU (relative imp.) | SARI (relative imp.) | Speedup |
|--------------------|----------------------|----------------------|---------|
| Best Seq2Seq model | 23.3 ( $\pm 0\%$ )   | 22.4 ( $\pm 0\%$ )   | -       |
| Emb2Emb            | <b>34.7</b> (+48.9%) | <b>25.4</b> (+13.4%) | 2.2×    |

Training models with a fixed-size bottleneck may be *easier, faster, and more effective* when training embedding-to-embedding!

## UNSUPERVISED TASK TRAINING

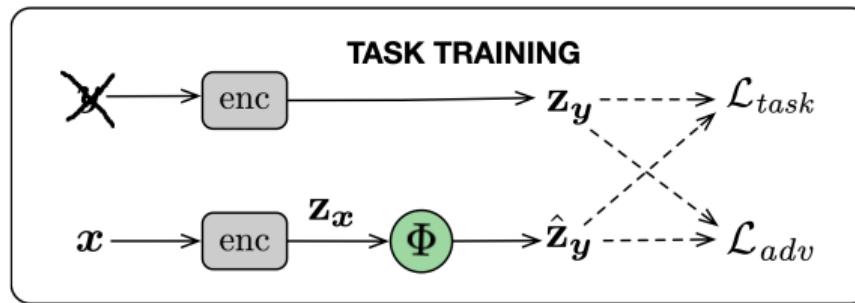
- Fixed-size bottleneck autoencoders are commonly used for unsupervised style transfer

## UNSUPERVISED TASK TRAINING

- Fixed-size bottleneck autoencoders are commonly used for unsupervised style transfer
- The goal is to change the **style** of a text, but retain the **content**: e.g., in machine translation, sentence simplification, sentiment transfer

## UNSUPERVISED TASK TRAINING

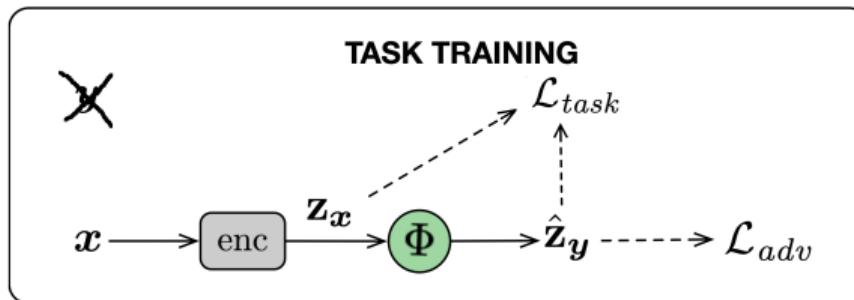
- Fixed-size bottleneck autoencoders are commonly used for unsupervised style transfer
- The goal is to change the **style** of a text, but retain the **content**: e.g., in machine translation, sentence simplification, sentiment transfer



- training objective:  $\mathcal{L} = \mathcal{L}_{task} + \lambda_{adv} \cdot \mathcal{L}_{adv}$

# UNSUPERVISED TASK TRAINING

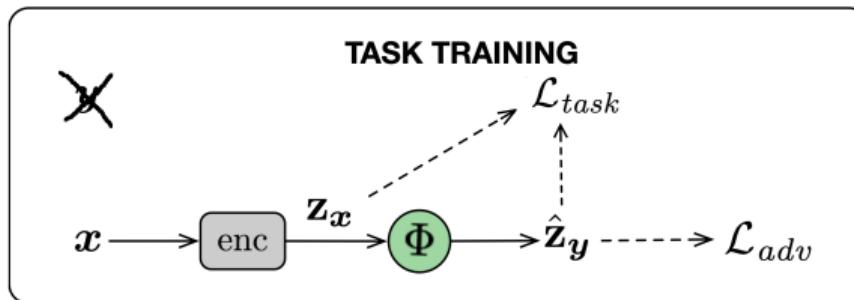
- Fixed-size bottleneck autoencoders are commonly used for unsupervised style transfer
- The goal is to change the **style** of a text, but retain the **content**: e.g., in machine translation, sentence simplification, sentiment transfer



- training objective:  $\mathcal{L} = \mathcal{L}_{task} + \lambda_{adv} \cdot \mathcal{L}_{adv}$
- $\mathcal{L}_{task}(\hat{\mathbf{z}}_y, \mathbf{z}_x) = \lambda_{sty}\mathcal{L}_{sty}(\hat{\mathbf{z}}_y) + (1 - \lambda_{sty})\mathcal{L}_{cont}(\hat{\mathbf{z}}_y, \mathbf{z}_x)$

# UNSUPERVISED TASK TRAINING

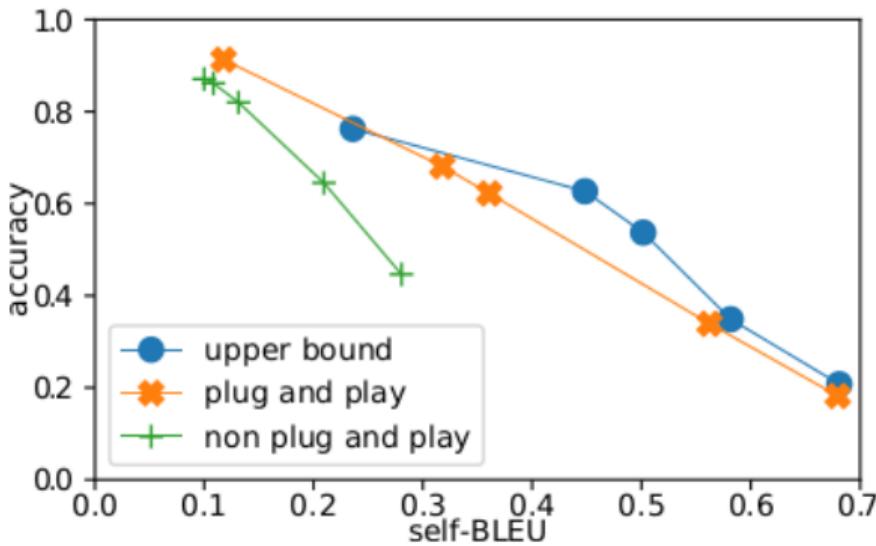
- Fixed-size bottleneck autoencoders are commonly used for unsupervised style transfer
- The goal is to change the **style** of a text, but retain the **content**: e.g., in machine translation, sentence simplification, sentiment transfer



- training objective:  $\mathcal{L} = \mathcal{L}_{task} + \lambda_{adv} \cdot \mathcal{L}_{adv}$
- $\mathcal{L}_{task}(\hat{\mathbf{z}}_y, \mathbf{z}_x) = \lambda_{sty}\mathcal{L}_{sty}(\hat{\mathbf{z}}_y) + (1 - \lambda_{sty})\mathcal{L}_{cont}(\hat{\mathbf{z}}_y, \mathbf{z}_x)$
- we set  $\mathcal{L}_{cont}$  to cosine distance, and  $\mathcal{L}_{sty}$  to a style classifier's negative log-likelihood of the target class

## UNSUPERVISED STYLE TRANSFER EXPERIMENTS

- Yelp sentiment transfer dataset: transform reviews with negative sentiment into reviews with positive sentiment (accuracy), but retain content (self-BLEU)
- if we have labels for only 10% of the data, how much better is a plug and play model?



### Effect of pretraining

By leveraging autoencoder pretraining on unlabeled data, our plug and play method offers a distinct advantage on unsupervised style transfer!

# CONCLUSION

In this talk...

# CONCLUSION

In this talk...

- we propose to learn in the embedding space of a pretrained autoencoder, training embedding-to-embedding (Emb2Emb).

# CONCLUSION

In this talk...

- we propose to learn in the embedding space of a pretrained autoencoder, training embedding-to-embedding (Emb2Emb).
- we discuss why it's important to keep the predicted embedding on the manifold of the autoencoder, and how to achieve that.

# CONCLUSION

In this talk...

- we propose to learn in the embedding space of a pretrained autoencoder, training embedding-to-embedding (Emb2Emb).
- we discuss why it's important to keep the predicted embedding on the manifold of the autoencoder, and how to achieve that.
- we demonstrate that a plug and play method like ours has a distinct advantage on unsupervised style transfer.

# CONCLUSION

In this talk...

- we propose to learn in the embedding space of a pretrained autoencoder, training embedding-to-embedding (Emb2Emb).
- we discuss why it's important to keep the predicted embedding on the manifold of the autoencoder, and how to achieve that.
- we demonstrate that a plug and play method like ours has a distinct advantage on unsupervised style transfer.

**Additionally,** our paper...

- presents an architecture for the mapping  $\Phi$  that is better than just MLPs.

# CONCLUSION

In this talk...

- we propose to learn in the embedding space of a pretrained autoencoder, training embedding-to-embedding (Emb2Emb).
- we discuss why it's important to keep the predicted embedding on the manifold of the autoencoder, and how to achieve that.
- we demonstrate that a plug and play method like ours has a distinct advantage on unsupervised style transfer.

**Additionally,** our paper...

- presents an architecture for the mapping  $\Phi$  that is better than just MLPs.
- demonstrates how to further improve the performance on unsupervised style transfer at inference time.

# *THANK YOU*